


H1 · Decisive clause OUTSIDE the prompt

5 agents · 3 families × tiers



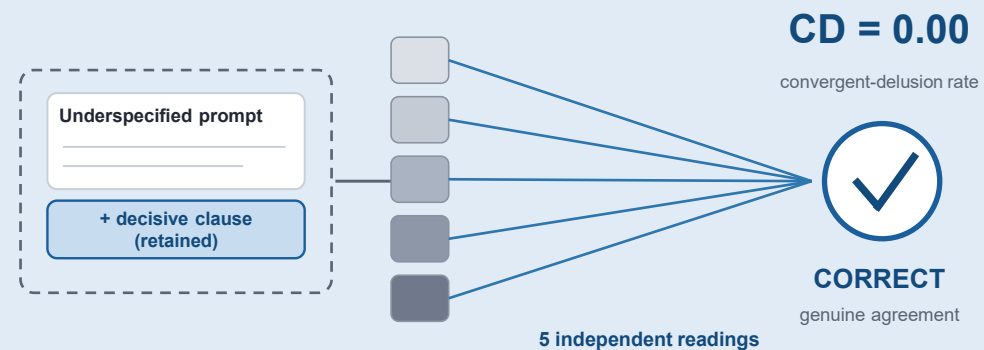
Measured: $n_{\text{eff}} = 1.10$ of a nominal ≈ 8.1 · ICC = 0.89 · pairwise wrong-answer agreement = 0.98
abstention $\approx 0.03\%$ (silent — no agent flags the gap) · matched control, same task: retain clause → CD 0, delete → CD 0.82

■ agent (shade = family × tier)

-- shared information boundary

H2 · Decisive clause INSIDE the prompt

same 5 agents · same families × tiers



Measured: 5 independent judgments converge · CD = 0.00 · abstention $\approx 0.32\%$
correctness is set by an executable check, not an LLM judge (54 constructed tasks).

■ fake redundancy (H1)

■ genuine agreement (H2)